

DocMind Dual LLM Gateway & Failover Specification

1. Primary Cloud Provider

The primary cloud inference provider is the Groq Cloud API utilizing the model qwen/qwen3.8-27b (configured as qwen3.8-27b). Groq provides rapid Time to First Token (TTFT) and an 8,192 token context window.

2. Local Fallback Daemon

DocMind maintains a secondary local Ollama instance running at <http://localhost:11434>.

Supported offline models include llama3.1:8b (4,096 token context window) and mistral:7b.

Gateway Routing & Circuit Breaking - Page 2

3. Automatic Failover Mechanism

When operating in 'auto' mode, DocMind transmits requests to Groq with a 2,500 ms connection timeout.

If Groq returns HTTP 429 (Rate Limit Exceeded), HTTP 5xx, or `AuthenticationError`, the gateway automatically fa

Failover occurs transparently without returning errors to the user.